

A Deterministic Schema Guardrail for LLM-Assisted SOC Alert Triage: A Negative-Result Investigation of Cross-Domain Classifier Transfer

Asma Imran^{1*}

^{1*}TECIP, Scuola Superiore Sant’Anna, Pisa, Italy.

Corresponding author(s). E-mail(s): asmaimran2003@gmail.com;

Abstract

LLM-based Security Operations Center (SOC) copilots increasingly rely on machine-learned guardrails to filter adversarial input before it reaches a triage model. We report a negative result from building and deploying such a guardrail in a real LangGraph-based alert-triage pipeline evaluated on Microsoft’s GUIDE dataset. Following a supervisor-suggested design (issue #10: regex fast-path → ML classifier → LLM route), we selected a TF-IDF + Logistic Regression detector — which outperformed Meta’s Llama Prompt Guard and Protect AI’s LLM Guard on a public jailbreak benchmark (F1 0.883 vs. 0.747 and 0.722, respectively, at three orders of magnitude lower latency) — and integrated it as a second guardrail stage scoped to the `AlertTitle` field. When evaluated against real GUIDE alert data rather than the benchmark distribution it was selected on, the classifier collapsed to near-chance performance (52.5% best-case accuracy across a threshold sweep; 1/20 true positives at the best-performing threshold). Root-cause analysis, rather than threshold tuning or retraining, revealed the actual cause: `AlertTitle` in the real dataset is a closed numeric-ID field (86,149 unique integer values), not free text — the GUIDE schema itself has no free-text column, so no amount of domain-matched training data could have closed the gap. We retired the classifier and replaced it with a deterministic type-validation guardrail that separates classes by construction rather than learned approximation, achieving 100% accuracy on a balanced synthetic set and zero false positives against 5,000 real `AlertTitle` values. We also report and fix a silent graph-wiring defect this change introduced — a dead-end node that caused every verdict to default to `FalsePositive` without raising an error — and the regression-test suite built to catch it. We argue that correctly diagnosing an unfalsifiable premise, and confirming a mismatch rather than deploying a broken classifier, is itself a reportable contribution for guardrail engineering.

Keywords: SOC triage, LLM guardrails, prompt injection, negative results, LangGraph, MITRE ATT&CK

1 Introduction

Motivation and context. Security Operations Centers face alert volumes that outpace human triage capacity, driving interest in LLM copilots that summarize, contextualize, and pre-classify alerts before an analyst sees them (Microsoft

Security Copilot [13]; Ferrag et al. [7]; Srinivas et al. [15]). A recurring design pattern in this space layers a fast, cheap guardrail (regex/keyword filtering) in front of a slower, more capable classifier, which in turn gates access to the LLM itself — mirroring the layered defense pattern used in general-purpose LLM guardrail frameworks such

as Meta’s LlamaFirewall [11] and Llama Prompt Guard [10], and Protect AI’s LLM Guard [14].

Problem definition. This layered pattern assumes the ML classifier stage has a text field to operate on, and that a classifier’s benchmark performance is representative of its deployment domain. We show both assumptions can silently fail in a real SOC pipeline: the field a classifier is scoped to may not be free text at all, and a classifier’s published F1 score may not transfer to the structured, terse text distribution of real alert data. Discovering this only after deployment — rather than before, from the schema — wastes the time-box allocated to a fix that cannot work.

Scope of this work. We report on a full design-build-evaluate-diagnose cycle carried out inside a working, evaluated LangGraph triage agent over the GUIDE dataset [12]: (1) selecting and integrating a second-stage ML guardrail per a specific architectural suggestion (issue #10), (2) evaluating it against real data rather than only the benchmark it was selected on, (3) finding it has no usable signal on that data, (4) root-causing *why* rather than retraining, and (5) replacing it with a deterministic guardrail that is provably correct for its target field, hard-gated into the live pipeline. We also report a real reliability defect this change introduced — a silently dead-ending graph edge — and the regression-test suite added to prevent recurrence.

Contributions.

1. **A negative-result methodology for guardrail transfer:** we show that a classifier chosen because it wins a public benchmark (F1 0.883 on a chat-style jailbreak eval set) can be statistically indistinguishable from chance on the actual deployment distribution, and we demonstrate that root-causing this from the data schema — rather than from threshold sweeps or retraining — is what actually resolves it, in one week instead of an open-ended tuning effort.
2. **A deterministic schema guardrail** for closed-vocabulary/numeric alert fields, which separates benign input from injection payloads by construction (type validity) rather than by a learned, thresholded score, and is safe to hard-gate for exactly that reason.
3. **A reliability-engineering finding and fix:** introducing the schema guardrail silently broke

the downstream classification routing (a dead-end node LangGraph does not flag at compile time), which a hardcoded default masked as a plausible-looking result instead of a crash. We document the defect, the fix, and the structural regression test (`tests/test_graph_wiring.py`) now guarding against recurrence.

Structure of this report. Section 2 surveys related work in LLM-based SOC copilots, guardrail frameworks, and MITRE ATT&CK-grounded LLM systems. Section 3 describes the triage pipeline, the MITRE ATT&CK retrieval stage, and the guardrail decision process in detail. Section 4 reports evaluation results exactly as committed in `experiments/results/`. Section 5 discusses the negative-result methodology and limitations. Section 6 concludes.

2 Related Work

Base table from docs/literature-review.md, extended with guardrail-framework and MITRE-grounded-LLM citations found for this paper. Author lists for the newly added arXiv entries still need a pass against publisher metadata before submission; see PROGRESS.md item R2.

LLM-based SOC copilots. Microsoft Security Copilot [13] demonstrates GPT-4 integrated into commercial SOC workflows, reporting reduced analyst triage time in a closed, proprietary system. Ferrag et al. [7] fine-tune and prompt LLMs for SIEM alert classification on CICIDS2017 and synthetic logs, showing LLMs can outperform rule-based classifiers — but on synthetic, not production-shaped, data. Wei et al.’s CORTEX [16] is the closest published architecture to ours: a multi-agent, evidence-grounded triage design that reduces false positives relative to single-model baselines, validating our own evidence-grounded (MITRE-context-fetching) LLM node design. Srinivas et al.’s survey [15] situates human-in-the-loop copilots as the dominant successful pattern in the field, over full automation — consistent with our pipeline’s human-review checkpoint for non-high-confidence verdicts. The GUIDE dataset itself originates from Microsoft’s own guided-response research [12], which we use directly as our evaluation corpus rather than a synthetic or network-flow substitute like CICIDS2017.

Guardrail frameworks for LLM systems. Meta’s LlamaFirewall [11] and Llama Prompt Guard [10], and Protect AI’s LLM Guard [14], are the current reference points for layered LLM input defense — regex/heuristic fast paths backed by learned classifiers for jailbreak and injection detection. All three are benchmarked, in our own prior work (PR #12), against a TF-IDF + Logistic Regression baseline on a chat-style jailbreak evaluation set, where the simpler classifier wins on both accuracy and latency. This paper extends that comparison past benchmark accuracy to deployment-distribution transfer, which none of [10, 11, 14] evaluate against structured, SOC-alert-shaped input as opposed to chat-style prompts.

MITRE ATT&CK-grounded LLM systems. Recent work grounds LLM threat-intelligence reasoning in MITRE ATT&CK via retrieval-augmented generation — e.g., structured retrieval over ATT&CK technique descriptions to reduce hallucination in attack-technique attribution [5], knowledge-graph-augmented retrieval over cybersecurity knowledge bases [4], and benchmarks for retrieval-augmented LLMs over heterogeneous cyber threat intelligence [3]. Our `fetch_mitre_context` pipeline stage (Section 3) follows the same retrieval-before-classification pattern established by Lewis et al.’s foundational RAG formulation [9], applied to MITRE technique context rather than open-domain text.

Prompt injection against structured/agent input. Most published prompt-injection taxonomies and defenses target free-text chat or agent-tool input [1, 2, 6]. Our contribution is a case study in the opposite direction: a field that looks like it should be defended as text (`AlertTitle`) is not text at all in the real schema, and the correct defense follows from the schema, not from the injection-detection literature.

3 Proposed Method

3.1 System architecture

The triage pipeline is a `LangGraph StateGraph` [8] (`src/agent/graph.py`) over a shared `AlertState`. Execution order:

Two guardrail stages run before any alert reaches the LLM or the MITRE-context retrieval step: a regex fast-path (`apply_regex_guardrail`,

pre-existing from Week 7 / issue #8) and, as of Week 9, a deterministic schema guardrail (`apply_schema_guardrail`) that replaced an ML classifier guardrail evaluated in Week 8. Routing after `fetch_mitre_context` (`route_by_context`) sends alerts with sufficient discriminative context fields to the LLM classifier and sparse alerts to a Random Forest fallback classifier (Week 6). Verdicts below high confidence are routed to a human-review checkpoint rather than ending directly — the same human-in-the-loop pattern surveyed by [15].

3.2 Data

GUIDE [12]: `datasets/GUIDE_train.csv` / `GUIDE_Test.csv`, Microsoft’s public real-world SOC alert/incident corpus, evaluated on real triage-labeled alerts throughout (not synthetic substitutes). `AlertTitle`, the field targeted by both guardrail designs in this paper, contains 86,149 unique values in the training split, all integers — confirmed by direct inspection, not inferred from the schema documentation alone.

3.3 Guardrail decision process — why the classifier was rejected, not deployed

Step 1 — issue-driven design. Issue #10 requested a second-stage ML classifier behind the regex fast-path, citing prior benchmark work comparing Meta Llama Prompt Guard and Protect AI LLM Guard [10, 11, 14]. We pulled that benchmark data and additionally evaluated a TF-IDF + Logistic Regression detector, which won on both axes that matter for a gate placed in front of every alert: F1 0.883 vs. 0.747 (Prompt Guard) and 0.722 (LLM Guard); median latency 0.79 ms vs. 179 ms and 183 ms — roughly three orders of magnitude faster (`experiments/results/` benchmark data referenced in PR #12).

Step 2 — integration and a real bug. The classifier was wired in as `src/agent/ml_guardrail.py`, scoped to `AlertTitle`. Cross-version scikit-learn unpickling (model trained under 1.7.1, environment running 1.7.2) initially produced non-deterministic scores; pinning `scikit-learn==1.7.1` (later committed to `requirements.txt` in PR #14, after an initial local-only fix was caught in review) resolved it.

```

START -> build_context -> regex_guardrail
      -> [continue] -> schema_guardrail
                    -> [continue] -> fetch_mitre_context
                                -> (classify_with_llm | rf_fallback)
                                -> parse_verdict -> (END | human_review)
                    -> [blocked] -> human_review
      -> [blocked] -> human_review

```

Fig. 1 Triage pipeline control flow (`src/agent/graph.py`). ASCII rendering ported as-is from the Markdown draft; a proper vector diagram is pending (PROGRESS.md item T2a).

Step 3 — deployment-distribution evaluation, not just benchmark evaluation. Run against real GUIDE alert titles, the detector flagged 100% of a 9-alert and a 30-alert batch as injection. A single benign sentence (“please schedule a team meeting for next Tuesday”) scored 0.726 against a 0.5 flag threshold. This prompted a purpose-built 40-row SOC-domain evaluation set (`experiments/soc_domain_eval_v1.csv`: 20 benign SOC alert titles, 20 injection attempts phrased as alert-field text), scored via `experiments/soc_domain_eval.py`.

Step 4 — a bug inside the negative result. Initial scoring (PR #12) had `score_text()` reading `classifier.predict_proba(X)[0][0]` — index 0, the benign-class probability — instead of `[0][1]`, the injection-class probability the guardrail was meant to threshold on. This inverted the *direction* of every score: the clearest injection example in the eval set (an HTML-comment-style override payload) scored lowest of all 40 rows (0.289), and a mundane benign alert scored highest (0.979). A reviewer (PR #14) caught the index bug; correcting it flipped both extremes as expected (the same injection example rose to 0.711, the same benign alert fell to 0.021) but did not change the conclusion — see Section 4.

Step 5 — root cause, not retraining. Before time-boxing a fine-tune of the classifier on SOC-domain-labeled text (the roadmap’s original Aug 9 plan), we checked what `AlertTitle` actually contains. It is a closed numeric-ID field, confirmed against 86,149 unique integer values in `GUIDE_train.csv`; the GUIDE paper’s own alert-dataframe schema lists no free-text categorical column at all. A TF-IDF classifier cannot find linguistic signal in a field that was never linguistic — no amount of domain-matched retraining data changes that. We retired the fine-tune plan on this basis rather than spending the time-box confirming a predictable failure.

Step 6 — the correct guardrail, and why it can be hard-gated. `src/agent/schema_guardrail.py` replaces the classifier with `validate_field_types()`: for each field in `EXPECTED_NUMERIC_FIELDS = {"AlertTitle", "DetectorId"}`, the value must be an int/float, or a string that parses cleanly as one; anything else — including any injection payload, which is definitionally not a valid integer — fails the check. This separates the two classes *by construction*, not by a learned, thresholded approximation, which is precisely why it is safe to hard-gate (`regex_guardrail` → `schema_guardrail` → `fetch_mitre_context`, `src/agent/graph.py`) in a way the probabilistic ML classifier never was: there is no accuracy/threshold tradeoff to reason about at deployment time.

3.4 A reliability defect introduced by the fix, and the regression test built to catch it

Wiring `schema_guardrail` into `graph.py` deleted the existing conditional edge from `fetch_mitre_context` to `(classify_with_llm | rf_fallback)` without restoring it. `LangGraph` does not validate dead-end nodes at compile time, so the graph compiled and ran to completion — but every alert silently stopped after the MITRE-context lookup, producing no `predicted_label`. This was invisible in practice because `evaluate.py` read the missing key with a default (`result.get("predicted_label", "FalsePositive")`) instead of raising, so every alert silently scored as a hardcoded guess rather than crashing or erroring visibly. `src/app.py`, `benchmark.py`, and `run_agent.py` all invoke the same compiled graph and were equally affected. No repository test previously existed to catch a structurally dead-ended graph.

The fix restores the deleted conditional edge. `tests/test_graph_wiring.py` now asserts, structurally, that every non-END node in the compiled graph has at least one outgoing edge, and separately asserts an end-to-end invocation on a sparse alert produces a real, non-default `predicted_label` — turning this specific bug class into a standing regression check rather than a one-time fix. `tests/test_schema_guardrail.py` and `tests/test_ml_guardrail.py` similarly convert the manual accuracy spot-checks from Sections 3.3–4 into assertions.

4 Evaluation

All figures below are read directly from the named JSON file — verify against the file before citing in the final version; do not hand-round.

4.1 Guardrail classifier selection benchmark (pre-deployment)

Source: PR #12 benchmark table (guardrail_comparison.json / eval_dataset_v2.csv, 1000-row balanced chat-jailbreak eval set — external to this repo, cited for classifier-selection justification only).

4.2 Deployment-distribution evaluation: SOC-domain eval set (post index-bug fix)

Source: `experiments/results/soc_domain_eval_results.json`, 40-row balanced synthetic set (20 benign / 20 injection), corrected [0][1] scoring.

Best achievable accuracy across the sweep is 52.5% (threshold 0.7), with only 1 of 20 true injections recovered even at that point — at chance for a balanced set. At threshold ≥ 0.8 the detector predicts “injection” for essentially no input, regardless of true label. This is presented as evidence of *no usable signal* on this domain, not a miscalibration recoverable by threshold tuning.

4.3 Schema guardrail evaluation (replacement)

Source: `experiments/results/schema_guardrail_eval.json`, generated by `experiments/schema_guardrail_eval.py`

(added for this paper — see `PROGRESS.md` item E1).

- Synthetic (20 benign numeric IDs vs. the 20 injection strings from the eval set above): **100% accuracy**, 0 false positives, 0 false negatives — reproduced independently of the PR #15 description.
- Real data (5,000 `AlertTitle` values sampled from `GUIDE_train.csv`, in row order — not randomly sampled): **0 false positives out of 5,000**, reproduced independently of the PR #15 description. `experiments/results/schema_guardrail_eval.json` now records `"status": "ran", "n_sampled": 5000, "n_false_positives": 0`. This closes item E1 — both halves of the guardrail’s headline claim now trace to a committed, reproducible artifact rather than PR prose.

4.4 Random Forest baseline

Source: `experiments/results/baseline_metrics.json` (`n_train=79580`, `n_test=19895`, `random_state=42`).

4.5 End-to-end agent evaluation, post graph-wiring fix

Source: `experiments/results/agent_metrics_post_graph_fix_week9.json` (sample size 30).

Predictions spread across all three classes post-fix (12 BenignPositive / 10 TruePositive / 8 FalsePositive; 9 alerts routed via LLM, 21 via RF fallback) — the qualitative signal that the graph is classifying again rather than collapsing to a single default. *Note for Section 5: this 30-sample post-fix run is smaller than the larger agent_metrics_real_v3.json 300-sample run reported elsewhere at accuracy 0.28 / macro F1 0.1836 against an RF baseline of 0.374/0.296 on that sample — these two files are not directly comparable (different sample sizes/dates) and reconciling which is the authoritative end-to-end number for the paper is an open item; see PROGRESS.md item E2.*

4.6 Scalability benchmark

Source: `experiments/results/week7_scalability_benchmark.json`.

Table 1 Pre-deployment classifier selection benchmark.

Detector	F1	Median latency	Throughput
TF-IDF + Logistic Regression	0.883	0.79 ms	1258 alerts/s
Meta Llama Prompt Guard	0.747	179 ms	5.6 alerts/s
Protect AI LLM Guard	0.722	183 ms	5.5 alerts/s

Table 2 SOC-domain threshold sweep, post index-bug fix.

Threshold	TP	FP	TN	FN	Accuracy	Precision	Recall	F1
0.5	3	3	17	17	0.500	0.500	0.15	0.231
0.6	1	1	19	19	0.500	0.500	0.05	0.091
0.7	1	0	20	19	0.525	1.000	0.05	0.095
0.8	0	0	20	20	0.500	0.000	0.00	0.000
0.9	0	0	20	20	0.500	0.000	0.00	0.000
0.95	0	0	20	20	0.500	0.000	0.00	0.000

- Regex guardrail microbenchmark: mean check time 3.616 μ s/alert over 10,000 iterations; 10,000/10,000 injection alerts blocked, 0/10,000 benign alerts blocked (at that stage’s test set).
- LLM path (1 worker, 30 prompts): throughput 0.5672 alerts/s, mean latency 1.7629 s — dominated by remote Groq inference and network latency, not local compute.

Full table — CPU/memory/worker-scaling rows — to be transcribed in full for the final version; only the headline rows are drafted here. See PROGRESS.md item E3.

5 Discussion

The negative-result methodology, and why it matters more than a working-but-wrong classifier. A classifier that appears to work — because it was evaluated only on the benchmark distribution it was selected against — is a worse outcome for a production guardrail than one that is honestly shown not to work on the deployment distribution. Section 4.1 shows the TF-IDF classifier winning decisively against two established guardrail frameworks; had that benchmark result alone been used to justify deployment, the guardrail would have shipped in a state indistinguishable from random on real input (Section 4.2), silently passing roughly half of actual injection attempts while also blocking legitimate alerts. The distinction argued in this paper is not

“domain mismatch exists” — a known risk in ML deployment generally — but that *diagnosing it correctly* (checking the schema before retraining) converts an open-ended, unbounded tuning problem into a one-week, closed investigation with a deterministic replacement. Section 3.3, Step 5 is the paper’s central methodological claim.

The graph-wiring defect as a reliability-engineering contribution. Section 3.4’s defect is notable less for its mechanism (a deleted graph edge) than for how it failed: silently, past a compiled graph, past a running pipeline, into a hardcoded default that looked like a real (if pessimistic) classification. No exception, log line, or crash surfaced it; it was only caught by an independent post-merge audit, not by the change’s own review. We argue this is a generalizable risk for any LangGraph-style pipeline where a downstream consumer applies a lenient default to a missing key — the failure mode is structurally identical to an ML guardrail that silently returns a plausible-but-wrong score, and the fix (a structural graph-completeness test, not just a functional one) generalizes past this specific bug.

Limitations

- The SOC-domain evaluation set (Section 4.2) is a 40-row synthetic set, not a labeled real-world injection corpus against structured alert fields — no such public corpus was identified during

Table 3 Random Forest baseline metrics.

Metric	Value
Accuracy	0.7718
Macro F1	0.7505
BenignPositive F1	0.8066
FalsePositive F1	0.6561
TruePositive F1	0.7889

Table 4 End-to-end agent evaluation, post graph-wiring fix ($n = 30$).

Metric	Value
Accuracy	0.5333
Macro F1	0.5337
RF baseline (same file, for comparison)	acc 0.7718, F1 0.7505

this work (see Related Work, Prompt injection paragraph).

- The end-to-end agent evaluation samples (30 and 300 rows, Section 4.5) are small relative to the GUIDE test split, and the two available result files are not reconciled to a single authoritative number — flagged as `PROGRESS.md` item E2, to resolve before submission.
- The real-data schema-guardrail sample (Section 4.3) is the first 5,000 rows of `GUIDE_train.csv` in file order, not a random sample — row order in the source file is not known to be free of systematic bias (e.g., chronological or org-clustered ordering), so this should be read as “no false positives in the first 5,000 rows encountered” rather than a claim about the full 86,149-value population.
- `EXPECTED_NUMERIC_FIELDS` includes `DetectorId` on the basis of the GUIDE paper’s schema description rather than the same direct large-sample inspection applied to `AlertTitle`; the module docstring notes this as confirmed but the inspection methodology is less thorough than `AlertTitle`’s.
- The guardrail-selection benchmark (Section 4.1) is drawn from an eval set built for a companion project’s chat-jailbreak use case, not this project’s SOC-alert domain — which is exactly the mismatch this paper reports, but it also means Section 4.1’s numbers cannot be used as evidence about SOC-domain performance, only as the justification for which classifier was worth testing further.

6 Conclusion

Draft — expand once Sections 4/5 are fully reconciled against committed result files (PROGRESS.md items E2, E3).

We built, evaluated, and rejected a second-stage ML guardrail for an LLM-assisted SOC triage pipeline after confirming — through direct inspection of the deployment data rather than threshold tuning — that its target field is not the kind of data the classifier can ever operate on. We replaced it with a deterministic schema guardrail that is correct by construction for that field and safe to hard-gate, unlike the probabilistic classifier it replaced. In fixing the schema guardrail’s integration, we found and fixed a silent graph-wiring defect that had been masking every verdict as a default classification, and built a structural regression test suite to prevent recurrence. We argue the central contribution is methodological: correctly diagnosing why a plan cannot work, before spending the time-box that would have confirmed it empirically, is itself a reportable and reusable result for guardrail engineering in structured-data domains.

Declarations

- Funding: Not applicable.
- Conflict of interest/Competing interests: Not applicable.
- Ethics approval: Not applicable.
- Consent to participate: Not applicable.

- Consent for publication: Not applicable.
- Availability of data and materials: GUIDE dataset [12], publicly available from Microsoft. Evaluation result files committed under `experiments/results/` in the project repository.
- Code availability: <https://github.com/AI-Security-Internships-2026/02-soc-copilot-threat-analysis>
- Authors' contributions: Not applicable (single author).

References

- [1] AUTHORS TO VERIFY (2024) Breaking down the defenses: A comparative survey of attacks on large language models. arXiv preprint arXiv:240304786 Authors pending verification, PROGRESS.md item R2
- [2] AUTHORS TO VERIFY (2024) Security and privacy challenges of large language models: A survey. arXiv preprint arXiv:240200888 Authors pending verification, PROGRESS.md item R2
- [3] AUTHORS TO VERIFY (2025) CTICConnect: A benchmark for retrieval-augmented LLMs over heterogeneous cyber threat intelligence. arXiv preprint arXiv:251011974 Authors pending verification, PROGRESS.md item R2
- [4] AUTHORS TO VERIFY (2025) KGA-agent4CTI: unlocking the power of LLM in threat intelligence. Cybersecurity <https://link.springer.com/article/10.1186/s42400-025-00458-2> – title/authors pending verification, PROGRESS.md item R2
- [5] AUTHORS TO VERIFY (2025) LLM-powered threat intelligence: a retrieval-augmented generation approach for cyber attack investigation. PeerJ Computer Science <https://peerj.com/articles/cs-3371/> – title/authors pending verification, PROGRESS.md item R2
- [6] AUTHORS TO VERIFY (2026) The landscape of prompt injection threats in LLM agents: From taxonomy to analysis. arXiv preprint arXiv:260210453 Authors pending verification, PROGRESS.md item R2
- [7] Ferrag MA, et al (2024) Automated cyber threat intelligence: LLMs for SIEM alert triage. arXiv preprint arXiv:240708628
- [8] LangChain (2024) LangGraph: Multi-agent workflows / agent orchestration framework. Software, released January 2024, <https://www.langchain.com/langgraph>
- [9] Lewis P, Perez E, Piktus A, et al (2020) Retrieval-augmented generation for knowledge-intensive NLP tasks. In: Advances in Neural Information Processing Systems (NeurIPS)
- [10] Meta AI (2024) Llama prompt guard — model card. https://github.com/meta-llama/PurpleLlama/blob/main/Prompt-Guard/MODEL_CARD.md
- [11] Meta AI (2025) LlamaFirewall: An open source guardrail system for building secure AI agents. arXiv preprint arXiv:250503574
- [12] Microsoft (2024) AI-driven guided response for security operation centers with microsoft copilot for security. arXiv preprint arXiv:240709017 GUIDE dataset
- [13] Microsoft Security Team (2024) Microsoft security copilot. Industry tool, <https://www.microsoft.com/en-us/security/business/ai-machine-learning/microsoft-copilot-for-security>
- [14] Protect AI (2023) LLM guard — the security toolkit for LLM interactions. Open-source software, <https://github.com/protectai/llm-guard>
- [15] Srinivas S, Kirk R, Zendejas E, et al (2025) AI-augmented SOC: A survey of LLMs and agents for security automation. Informatics 5(4):95
- [16] Wei T, Tay Y, Liu X, et al (2025) CORTEX: Collaborative LLM agents for high-stakes alert triage. arXiv preprint arXiv:251000311